

# When Top-down Meets Bottom-up: Detecting and Exploiting Use-After-Cleanup Bugs in Linux Kernel

Lin Ma\*, Duoming Zhou\*, Hanjie Wu<sup>†</sup>, Yajin Zhou\*, Rui Chang\*,  
Hao Xiong\*, Lei Wu\*, Kui Ren\*

[\\*{linma, duoming, yajin zhou, crix1021, mart1n, lei wu, kuiren}@zju.edu.cn](mailto:{linma, duoming, yajin zhou, crix1021, mart1n, lei wu, kuiren}@zju.edu.cn), Zhejiang University

[†hanjiew@andrew.cmu.edu](mailto:†hanjiew@andrew.cmu.edu), Carnegie Mellon University



# Kernel Vulnerability Mining

---

- **From user-space side (Syscall)**

- [syzkaller\[1\]](#), [Difuze\[CCS '17\]](#), [Healer\[SOSP '21\]](#) ...

- From device side (Interrupt, DMA)

- syzkaller-external ([usb/bt/nfc/wireless\[2\]](#)), [PrintFuzz\[ISSTA '22\]](#) ...


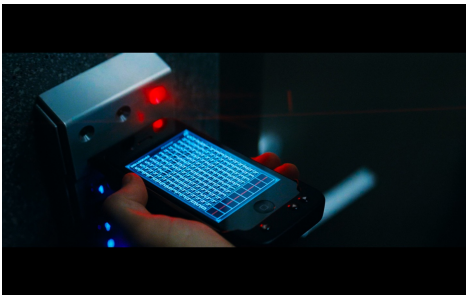
[1] <https://github.com/google/syzkaller>

[2] [https://github.com/google/syzkaller/blob/master/docs/linux/wifi\\_fuzzing.md](https://github.com/google/syzkaller/blob/master/docs/linux/wifi_fuzzing.md)



# Kernel Vulnerability Mining

- From user-space side (Syscall)
  - [syzkaller\[1\]](#), [Difuze\[CCS '17\]](#), [Healer\[SOSP '21\]](#) ...
- From device side (Interrupt, DMA)
  - syzkaller-external ([usb/bt/nfc/wireless\[2\]](#)), [PrintFuzz\[ISSTA '22\]](#) ...



# Kernel Vulnerability Mining

---


- From user-space side (Syscall)

- [syzkaller\[1\]](#), [Difuze\[CCS '17\]](#), [Healer\[SOSP '21\]](#) ...

- From device side (Interrupt, DMA)

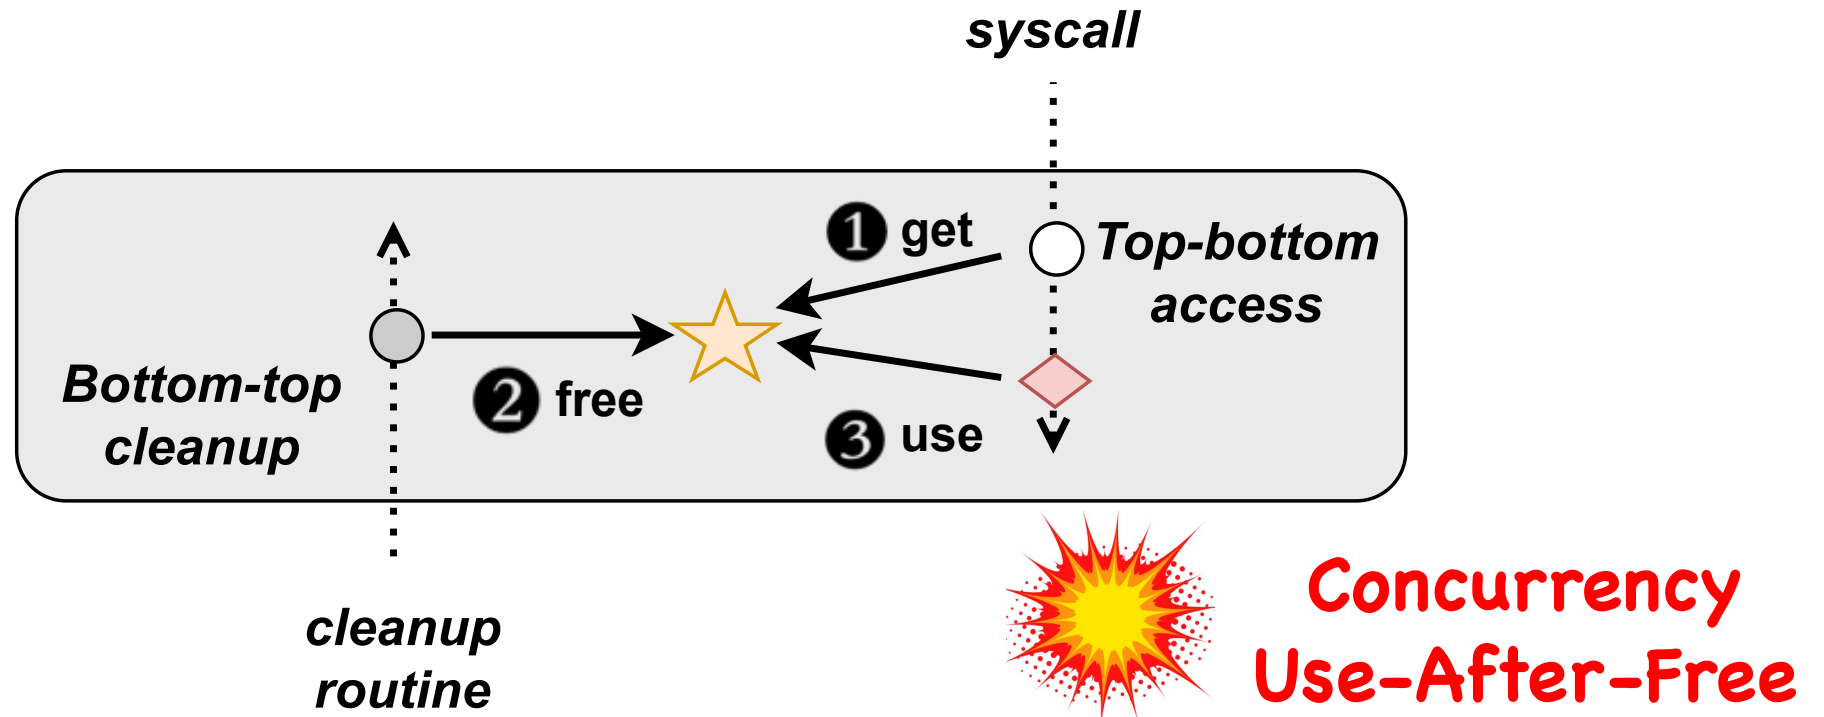
- [syzkaller-external \(usb/bt/nfc/wireless\[2\]\)](#), [PrintFuzz\[ISSTA '22\]](#) ...

***somehow  
combine?***



# UAC (Use-After-Cleanup)

When a device is **detached**, the running syscall may be unaware and **attempt to use** an already **released** object



# However

---

The community does not pay much attention to UAC bug, which **was** improperly treated as a low-security impact

- Concurrency issue seems need some luck to trigger 🙄
- Detachment seems need real hardware (or root-privilege-required virtual device) ⚙️

# However again, CVE-2021-3573

---

- ~~Concurrency issue seems need luck to trigger~~

Stably triggered 👍

- ~~Detachment seems need real hardware~~

From low-privilege user-space 👍

and Achieve Local Privilege Escalation (LPE) Attack

check [demo\[1\]](#) here and [details\[2\]](#) here

Blue  
Klotski



[1] <https://github.com/uacatcher/uacatcher-repo/blob/main/demo.gif>

[2] <https://f0rm2l1n.github.io/2021-07-23-Blue-Klotski/>

# Motivations

---

1. UAC bug is exploitable and with high risks  
*which is different from the widely thought*
2. UAC bug cannot be systematically detected by existing tools
  - Not for race at all ([deadline\[SP '18\]](#), [DASC\[ATC '18\]](#), ...)
  - Just for data race ([Razzer\[SP '19\]](#), [Krace\[SP '20\]](#), ...)
  - For CUAF but not concern cleanup ([DCUAF\[ATC '19\]](#))




# Our Solution: UACatcher

---

- Detect Use-After-Cleanup Bug via **Static Analysis**
- and Estimate **Exploitable** Bugs from detected ones

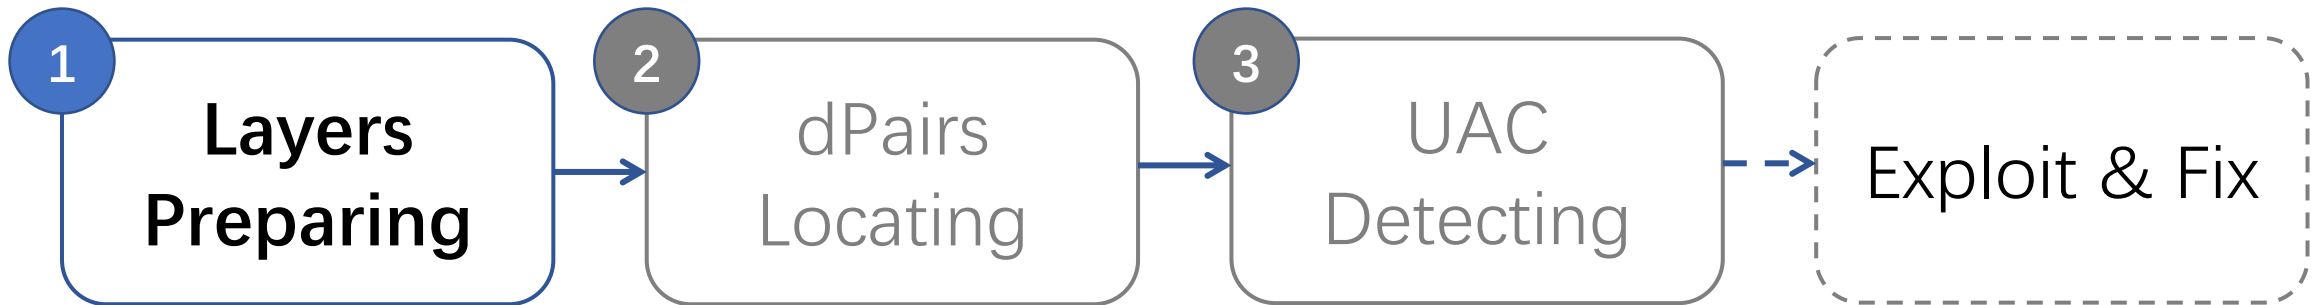
## Three Main Phases Overview



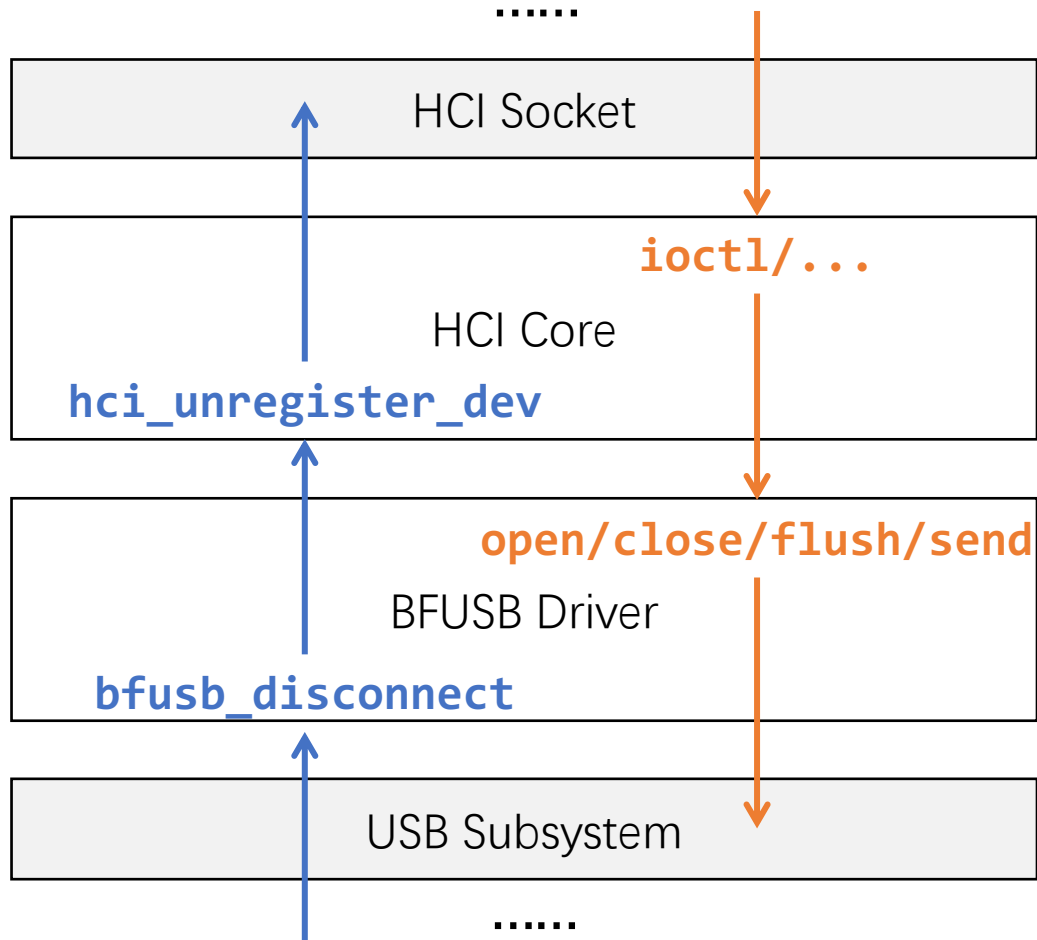
# 1 Layers Preparing

---

- Analyzing the entire kernel is too **expensive** for static analysis
  - the kernel follows a **layered architecture**
  - routines can be **split** into **layers granularity**



# Example: BT



- **unreg-entry functions**
  - cleanup routine
- **interface functions**
  - syscall routine
- Collectively referred to as **boundary functions**
- Found via **model** and **heuristics** methods

# Example: BT - Find driver unreg-entry functions

```
struct device_driver {  
    ...  
    int (*remove) (struct device* dev)  
    ...  
};
```

```
struct usbdrv_wrap {  
    ...  
    struct device_driver driver;  
    ...  
};
```

```
struct usb_driver {  
    ...  
    void (*disconnect) (...)  
    ...  
    struct usbdrv_wrap drvwrap;  
    ...  
}
```

1. find the driver type (`usb_driver`)
2. annotate the pointer field for handling cleanup (`disconnect`)
3. derive the actual unreg-entry function from specific driver struct (`bfusb_disconnect` from `bfusb_driver`)

```
static struct usb_driver bfusb_driver =  
{  
    .disconnect    = bfusb_disconnect,  
    ...  
}
```



# Example: other BT boundary functions

---

With the driver layer unreg-entry pinpointed

- Find the above layers unreg-entry from bottom-up via **characteristics-based heuristics**
  - `void hci_unregister_dev(struct hci_dev *hdev)`
- Find the interface functions from top-down as **packed code pointers [1]**
  - `open/close/flush/send` from `struct hci_dev`


[1] Zhang, Tong, et al. "Pex: A permission check analysis framework for linux kernel." *28th USENIX Security Symposium*. 2019.



## 2 dPairs Locating

---

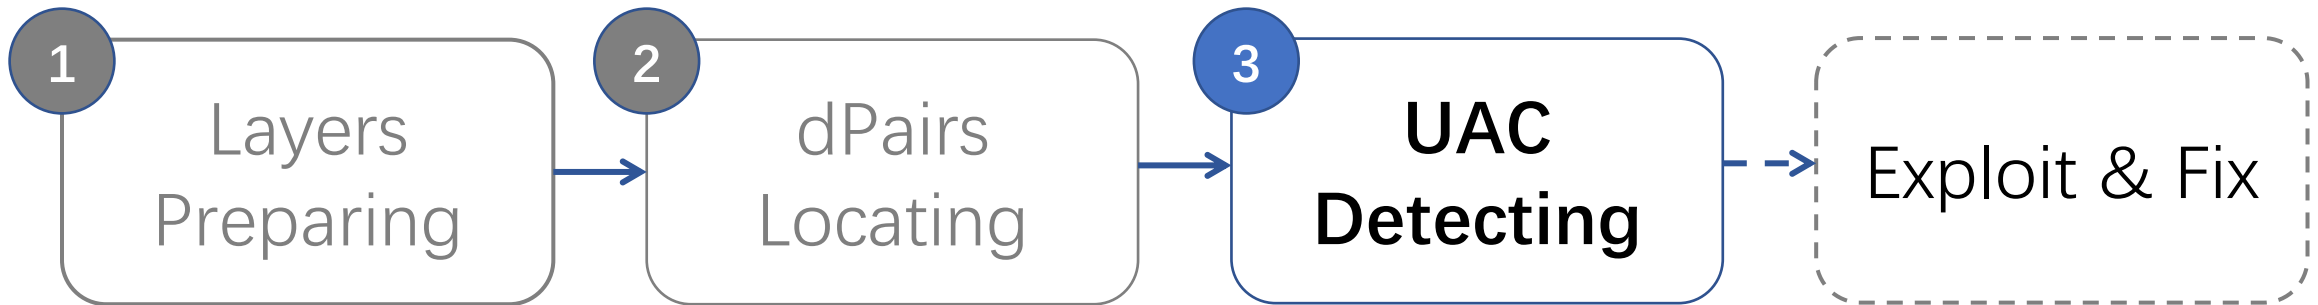
- Locate all **deallocation** & **dereference** site pairs (dPairs)
  - Track all call sites to dealloc functions (e.g., kfree)
  - Leverage points-to analysis
  - *Add filters to discard false result*



### 3 UAC Detecting

---

- Detect if a located dPair causes UAC or not
  - *routine-switch-point algorithm*
- Also **Estimate** detected bugs



# Example: UAC-free code

---

## cleanup routine

```
1 mutex_lock(&ndev->req_lock)
2 if !test_and_clear_bit(NCI_UP, &ndev->flags)
    return
...
3 mutex_unlock(&ndev->req_lock)
4 destroy_workqueue(ndev->cmd_wq)
   // deallocation
...
```

## syscall routine

```
5 mutex_lock(&ndev->req_lock)
6 if !test_bit(NCI_UP, &ndev->flags)
    goto 8
...
7 queue_work(ndev->cmd_wq, ...)
   // dereference
...
8 mutex_unlock(&ndev->req_lock)
```

Use **synchronizations** and **path constraints** to make sure

- *dereference would not happen after the deallocation*





# Routine-switch-point Algorithm

---

Find the routine-switch point that

- schedule dereference after the deallocation

Forward & Backward to get all path constraints  $P$

- If  $P$  is satisfy-able, the dPair cause true UAC
- otherwise, safe one

check more details in the paper ☺



# UAC Bug Estimation

---

## 1. Race Window **Identification** and **Measurement**

concern on **time-consuming functions** (memory allocation/IO/...) and **time-controllable functions** (copy\_{from/to}\_user)

## 2. User-space Device Emulation

by **Pseudoterminal Device** and **Line Discipline**

support BT/NFC/HAMRADIO/CAN devices for now



# Implementation & Evaluation Setup

---


- CodeQL Queries
  - port from user-space to Linux kernel
- Python Code
  - driver for running QL and fetching result
  - CallGraph, CFG
  - routine-switch point algoirhm
- Evaluated with Linux kernel 5.11 (7289e26f395b)



# Evaluation Result (for each phase)


---

- **Layers Preparing**
  - **88** driver types, **1,856** layers
- **dPairs Locating Phase**
  - **136,628** dPairs located
  - **51,270** after filtering
- **UAC Detecting Phase**
  - **436** detected UAC
  - **346** true bugs
  - **277** been confirmed
  - **15** CVEs assigned



# Evaluation Result (compare with DCUAF<sup>1</sup>)

- Locate more dPairs: **51,270** > **16,832**
- Detect more real UAC bugs: **346** > **130**
- Get less False Positives: **20.6%** < **98.6%**



<sup>1</sup>Not open source, we developed a QL version of it as a comparison



# Evaluation Result (estimation)

## 13 exploitable UAC bugs (one from 5.15 kernel)

layer	description	impact
BT	Read connection information from a destroyed device	□
	Read authentication information from a destroyed device	□
	Add address to the blacklist of a destroyed device	◇
	Remove address from the blacklist of a destroyed device	◇
	Enqueue a work to a destroyed workqueue	★
	Fetch an already reclaimed connection object	★
	Fetch and dereference a NULL resource pointer	△
NFC	Enqueue a work to a destroyed workqueue	★
	Fetch and dereference a NULL resource pointer	△
AX25	Read address information from a destroyed device	□
	Copy controllable data into reclaimed buffer	★
	Fetch and dereference a NULL resource pointer	△
MCTP	Read route information from a destroyed device	□

- read from a destroyed object
- ◇ write to a destroyed object
- △ cause Denial of Service
- ★ complex task

# Conclusions

---

- UAC bugs are **high risks**
- UACatcher is the **first** tool that **systematically** detects UAC bugs in Linux kernel
- **346** true bugs found, **277** fixed, **15** CVEs assigned
- **13** exploitable UAC estimated



# Thank you!

---

Lin MA, ZheJiang University  
[linma@zju.edu.cn](mailto:linma@zju.edu.cn)

